

# CSC2100 Data Structures

## Hashing

Irwin King

[king@cse.cuhk.edu.hk](mailto:king@cse.cuhk.edu.hk)  
<http://www.cse.cuhk.edu.hk/~king>

Department of Computer Science & Engineering  
The Chinese University of Hong Kong



# Introduction

- Hashing is a technique used for performing **insertions**, **deletions** and **finds** in constant **average** time.
- Tree operations that require any ordering information among the elements are not supported efficiently.
  - See several methods of implementing the hash table.
  - Compare these methods analytically.
  - Show numerous applications of hashing.
  - Compare hash tables with binary search trees.



# Rectangular Arrays

		Column			
		1	2	3	4
Row	1				
	2				
	3				

row 1

row 2

row 3

row  $u_1$

col 1	col 2	...	col $u_2$
X	X	...	X
X	X	...	X
X	X	...	X
	:		
X	X	...	X

(a)



# Row- and Column-Major Ordering

- How does one index an entry in an array?
- Entry  $[i,j]$  goes to position  $ni + j$  for row-major ordering and  $i + jm$  for column-major ordering when the rows are numbered from 0 to  $m-1$  and the columns from 0 to  $n-1$  and entry  $[0,0]$  is at position 0.



# Row- and Column-Major Example

Row-Major:



Column-Major:



# Example


$a$	address	$a$
$a[1,1]$	$Alpha$	$a[1,1]$
$a[1,2]$	$Alpha + 1$	$a[2,1]$
$a[1,3]$	$Alpha + 2$	$a[3,1]$
$a[1,4]$	$Alpha + 3$	$a[1,2]$
$a[2,1]$	$Alpha + 4$	$a[2,2]$
$a[2,2]$	$Alpha + 5$	$a[3,2]$
• •	•	• •
• •	•	• •
• •	•	• •
$a[3,4]$	$Alpha + 11$	$a[3,4]$

Row-Major




Column-Major




# Implementation Example



With the  
corners  
identified:




# More Implementation Example






# 3D Array Implementation



# 3-D Array Implementation



# Problem

- In some applications the full use of the whole array is seldom.
- This leads to sparse array or matrix representation.
- For example, population count in a 2-D grid map.



# An Access Table

- One method to eliminate the multiplications needed in calculating the index to an entry is to use an access table.
- The array will contain the values  $0, n, 2n, 3n, \dots, (m-1)n$ .
- Then for all references to the rectangular array, the index for  $[i,j]$  is calculated by taking the entry in position  $i$  of the auxiliary table, adding  $j$ , and going to the resulting position.
- Again we see a trade-off between space used and execution speed.



# Example

		Columns				
		1	2	3	4	5
Rows	1		12		14	
	2					
	3		26		40	
	4					
	5	17		31		
	6		9			85



# Example

	<i>T</i>	Columns				
		1	2	3	4	5
R o w s	1		12		14	
	2					
	3		26		40	
	4					
	5	17		31		
	6		9			85

		Columns									
		1	2	3	4	5	6	7	8	9	10
	1	5	17								
	2										
	3										
	4										
	5										
	6										



# Example

Columns	
1	2
5	17

Columns	
1	2
1	12
3	26
6	9

Columns	
1	2
5	31

Columns	
1	2
1	14
3	40

Columns	
1	2
6	85

T	Columns				
	1	2	3	4	5
R o w s		12		14	
		26		40	
	17		31		
		9			85

Value	17	12	26	9	31	14	40	85
Row	5	1	3	6	5	1	3	6
	1	2			5	6		8
								9

FirstInCol	1	2	5	6	8	9
	1	2	3	4	5	6



# Tables: A New Abstract Data Type

- Functions: a **function** is defined in terms of two sets and a correspondence from elements of the first set to elements of the second.
- If  $f$  is a function from a set  $A$  to a set  $B$ , then  $f$  assigns to each element of  $A$  a unique elements of  $B$ .
- The set  $A$  is called the domain of  $f$ , and the set  $B$  is called the codomain of  $f$ . The subset of  $B$  containing just those element that occur as values of  $f$  is called the range of  $f$ .





# Example

- For a table, we call the domain the **index** set, and we call the codomain the **base type** or **value type**.
- For example, to index into the cell  $[2,3]$  the offset value may be 13 if the matrix size is  $[10,10]$ .



# An Abstract Data Type

- A **table** with index set  $I$  and base type  $T$  is a function from  $I$  into  $T$  together with the following operations.
  - Table access: Evaluate the function at any index in  $I$ .
  - Table assignment: Modify the function by changing its value at a specified index in  $I$  to the new value specified in the assignment.



# An Abstract Data Type

- Insertion: Adjoin a new element  $x$  to the index set  $I$  and define a corresponding value of the function at  $x$ .
- Deletion: Delete an element  $x$  from the index set  $I$  and restrict the function to the resulting smaller domain.



# Why Hash Table?

- Often, array indices are not natural identifiers for items that are to be stored, accessed, and retrieved.
- For example, let's try to store the list in an array.

beef	bellpepper	blackpepper	dillweed
onion	potato	olive	salt
cumin	carrot	mushroom	tomatopaste



# Problem

- While it is true that STORE and RETRIEVE are  $O(1)$  operations for arrays, that is only so if the indices are **known** and the value in the target of a STORE can be discarded.
- Without a complete set in hand it cannot be known that potato has index 10 in the sorted list of items.



# Solution

- Use item as a **KEY**--Because an index integer is not known on the entry of one of the items, it would be helpful if the item itself could be used as a key to index the cell where it will be stored.



# Solution

- A solution would be to convert the keys (the list items, here) into **unique** integers and use them as array indices.
- A function that does so is called a **hash function**.
- The conversion process is called **hashing**
- The storage structure is called a **hash table** or **scatter-storage**.



# Example Solution

- We may sum up the ASCII value from each character in the key, e.g.,  $a = 1, b=2, \dots, z = 26$ , so  $\text{beef} = 2+5+5+6=18$ .

Item	HF1 (Item)	Item	HF1 (Item)
beef	18	carrot	75
onion	67	salt	52
cumin	60	blackpepper	105
dillweed	74	olive	63
bellpepper	107	tomatopaste	145
potato	87	mushroom	122





# Introduction to Hashing

- Hashing is a technique used for performing insertions, deletions and finds in **constant** average time.
- Tree operations that require any **ordering** information among the elements are not supported efficiently.
  - See several methods of implementing the hash table.
  - Compare these methods analytically.
  - Show numerous applications of hashing.
  - Compare hash tables with binary search trees.



# General Idea

- Hash table data structure is merely an array of some fixed size, containing the keys.
- A key is a string with an associated value (for instance, salary information).
- Each key is mapped into some number in the range 0 to  $H\_SIZE - 1$  and placed in the appropriate cell.

$f : \text{key} \rightarrow \text{value}$



# General Idea

- The mapping is called a **hash function**, which ideally should be **simple** to compute and should ensure that any two distinct keys get **different** cells.
- This is difficult to achieve in reality since there are a finite number of cells and a virtually inexhaustible supply of keys.
- We seek a hash function that distributes the keys **evenly** among the cells.



# Issues

- Choosing the hashing function
  - How to make sure that one has selected a good function for the application
- Collision handling
  - How to handle conflict when two keys have the same location
- Deletion handling
  - How to deal with the table when items are being removed



# Example

0	
1	
2	
3	john 25000
4	phil 31250
5	
6	dave 27500
7	mary 28200
8	
9	



# Hash Tables

- We can continue to exploit table lookup even in situations where the key is no longer an index that can be used directly as in array indexing.
- What we can do is to set up a **one-to-one** correspondence between the keys by which we wish to retrieve information and indices that we can use to access an array.



# Hash Tables

- The idea of a **hash table** is to allow many of the different possible keys that might occur to be mapped to the same location in an array under the action of the index function.
- Others have called **scatter-storage** or **key-transformation**.



# Hash Function

- A **hash function** that takes a key and maps it to some index in the array.
- Often, two records may want to go to the same location.
- Therefore, a **collision** may occur and a **collision procedure** must be devised to handle this.





# Choosing a Hash Function

- The two principal criteria in selecting a hash function are that
  - it should be **easy** and **quick** to compute and that
  - it should achieve an **even distribution** of the keys that actually occur across the range of indices.



# Hash Function

- If the input keys are integers, then simply returning  $\text{key} \bmod H\_SIZE$  is generally a reasonable strategy.
- For example,  $\text{student ID} \bmod 10000$  would be a reasonable strategy.
- It is usually a good idea to ensure that the table size is **prime**.
- When the input keys are random integers, then this function is **simple** to compute and also distributes the keys **evenly**.



# A Simple Hash Function

INDEX

```
hash( char *key, unsigned int H_SIZE )  
{  
    unsigned int hash_val = 0;  
  
    /*1*/      while( *key != '\0' )  
  
    /*2*/      hash_val += *key++;  
  
    /*3*/      return( hash_val % H_SIZE );  
}
```



# Another Hash Function

INDEX

```
hash( char *key, unsigned int H_SIZE )  
  
{  
  
    return ( ( key[0] + 27*key[1] + 729*key[2] ) %  
H_SIZE );  
  
}
```



# Notes

- Assuming key has at least two characters plus the NULL terminator.
- 27 represents the number of letters in the English alphabet, plus the blank.
- 729 is  $27^2$ .
- This function only examines the first three characters, but if these are random, and the table size is 10,007, as before, then we would expect a reasonably equitable distribution.



# Quick Analysis

- Unfortunately, English is not random.
- Although there are  $26 \times 26 \times 26 = 17,576$  possible combinations of three characters (ignoring blanks), a check of a reasonably large on-line dictionary reveals that the number of different combinations is actually only **2,851**.
- Even if none of these combinations collide, only 28% of the table can actually be hashed to.



# A Good Hash Function

INDEX

```
hash( char *key, unsigned int H_SIZE )  
  
{  
  
    unsigned int hash_val = 0;  
  
    /*1*/      while( *key != '\0' )  
  
    /*2*/      hash_val = ( hash_val << 5 ) + *key++;  
  
    /*3*/      return( hash_val % H_SIZE );  
  
}
```



# Notes

- This hash function involves all characters in the key .
- It computes  $\sum_{i=0}^{Key\_Size-1} Key[Key\_Size - i] \cdot 32^i$
- The code computes a polynomial function (of 32) by use of Horner's rule.
- For instance, another way of computing  $h_k = k_1 + 27 k_2 + 27^2 k_3$  is by the formula  $h_k = ((k_3) * 27 + k_2) * 27 + k_1$ .
- Horner's rule extends this to an nth degree polynomial.





# Notes

- It is common to not use all the characters.
- The length and properties of the keys would influence the choice.
- The hash function might include a couple of characters from each field.



# Truncation

- Ignore part of the key, and use the remaining part directly as the index (considering non-numeric fields as their numerical codes).
- Example: If the keys are eight-digit integers and the hash table has 1000 locations, then the first, second, and fifth digits from the right make the hash function, so that 62538194 maps to 394.
- Truncation is a very fast method, but it often fails to distribute the keys evenly through the table.



# Folding

- Partition the key into several parts and combine the parts in a convenient way (often using addition or multiplication) to obtain the index.
- For example, 62538194 maps to  $625+381+94 = 1100$ , which is then truncated to 100.



# Modular Arithmetic

- Convert the key to an integer (using the preceding devices as desired), divide by the size of the index range, and take the remainder as the result.
- For example, 'abcd' =  $64+65+66+67 \bmod 100 = 62$ .



# Collision Resolution

- Open Hashing (Separate Chaining)
- Closed Hashing (Open Addressing)
  - Linear probing
  - Quadratic probing
  - Double hashing
- Rehashing




# Open Hashing

- The first strategy, commonly known as either **open hashing**, or **separate chaining**, is to keep a **list** of all elements that hash to the same value.
- We assume for this section that the keys are the first 10 perfect squares and that the hashing function is simply  $\text{hash}(x) = x \bmod 10$ . (The table size is not prime, but is used here for simplicity.)



# Open Hashing Example



# Find in Open Hashing

- Find
  - We use the hash function to determine which list to traverse.
  - We then traverse this list in the normal manner, returning the position where the item is found.





# Insert in Open Hashing

- Insert
- we traverse down the appropriate list to check whether the element is already in place.
- If the element turns out to be new, it is inserted either at the front of the list or at the end of the list.
- Sometimes new elements are inserted at the **front** of the list, since it is convenient and also because frequently it happens that recently inserted elements are the most likely to be accessed in the near future.



# Deletion in Open Hashing

- Deletion
  - The deletion routine is a straightforward implementation of deletion in a linked list.
  - First perform a FIND operation and then perform a delete operation of an item in a linked list.



# Advantages of Linked Storage

- Considerable space may be saved.
- It allows simple and efficient collision handling.
- It is no longer necessary that the size of the hash table exceed the number of records.
- Deletion becomes a quick and easy task in a chained hash table.



# Disadvantage of Linked Storage

- All the links require space.
- If the records are small this space usage is large when compared with the records.



# Closed Hashing (Open Addressing)

- Open hashing has the disadvantage of requiring pointers.
- This tends to slow the algorithm:
  - The time required to allocate new cells.
  - It requires the implementation of a second data structure.
- In a **closed hashing** system, if a collision occurs, **alternate cells** are tried until an empty cell is found.



# Closed Hashing

- For example, cells  $h_0(x)$ ,  $h_1(x)$ ,  $h_2(x)$ , ... are tried in succession where  $h_i(x) = (\text{hash}(x) + f(i)) \bmod H\_SIZE$ , with  $f(0) = 0$ .
- The function,  $f$ , is the **collision resolution** strategy.
- Because all the data goes inside the table, a bigger table is needed for closed hashing than for open hashing.
- Generally, the load factor should be below  $\alpha = 0.5$  for closed hashing.



# Insertion Operation Outline

- An array must be declared that will hold the hash table.
- Initializing all locations in the array to show that they are empty.
- To insert a record into the hash table, the hash function for the key is first calculated.
- If the corresponding location is **empty**, then the record can be inserted, or else
- if the keys are **equal**, then insertion of the new record would not be allowed. In this case, it becomes necessary to resolve the collision.



# Find Operation Outline

- To retrieve the record with a given key is entirely similar. The hash function for the key is computed.
- If the desired record is in the corresponding location, then the retrieval has succeeded;
- otherwise,
- while the location is **nonempty** and not all locations have been examined, follow the same steps used for collision resolution.
- If an **empty** position is found, or  $h_0$  have been considered, then no record with the given key is in the table, and the search is unsuccessful.





# Linear Probing

- The simplest method to resolve a collision is to start with the hash address (the location where the collision occurred) and do a **sequential search** for the desired key or an empty location.
- The problem with the above method is that the data become **clustered**:
- Records start to appear in long strings of adjacent positions with gaps between the strings.



# Example

	Empty Table	After 89	After 18	After 49	After 58	After 69
0				49	49	49
1					58	58
2						69
3						
4						
5						
6						
7						
8			18	18	18	18
9		89	89	89	89	89



# Problems

- The time to search for an empty cell may be long.
- The problem of **primary clustering** is essentially one of **instability**.
- If a few keys happen randomly to be near each other, then it becomes more and more likely that other keys will join in the cluster.
- Furthermore, the distribution will become progressively more unbalanced.



# Analysis

- It can be shown that the expected number of probes using linear probing is roughly

- Insertions and unsuccessful searches

$$\frac{1}{2}(1 + 1/(1 - \lambda)^2)$$

- Successful searched

$$\frac{1}{2}(1 + 1/(1 - \lambda))$$

- $\lambda$ , of a hash table is the ratio of the number of elements in the hash table to the table size.




# Analysis

- We assume a very large table and that each probe is independent of the previous probes.
- The expected number of probes in an unsuccessful search.
- The number of probes for a successful search = the number of probes required when the particular element was inserted.
- When an element is inserted, it is done as a result of an unsuccessful search.
- We can use the cost of an unsuccessful search to compute the average cost of a successful search.
- Since the fraction of empty cells is  $1 - \lambda$ , the number of cells we expect to probe is  $1/(1 - \lambda)$ .



# Probes vs. Load Factor



- **Dashed curves**-- linear probing
- **Solid curves**-- random collision resolution
- **S**-successful
- **U**-unsuccessful
- **I**-insert
- What it is saying is that the linear probing is not a very good method to handle collision.



# Notes

- If  $\lambda = 0.75$ , then the formula above indicates that 8.5 probes are expected for an insertion in linear probing.
- If  $\lambda = 0.9$ , then 50 probes are expected.
- This compares with 4 and 10 probes for the respective load factors if clustering were not a problem.
- We see from these formulas that linear probing can be a bad idea if the table is expected to be more than half full.
- If  $\lambda = 0.5$ , however, only 2.5 probes for insertion and only 1.5 probes are required for a successful search.



# Quadratic Probing

- Quadratic probing avoid the primary clustering problem of linear probing.
- If there is a collision at hash address  $H$ , the method call **quadratic probing** looks in the table at locations  $h+1$ ,  $h+4$ ,  $h+9$ , ..., that is, at locations  $h + i^2 \pmod{\text{hashsize}}$  for  $i=1, 2, \dots$
- This reduces clustering, but it is not obvious that it will probe all locations in the table, and in fact it does not.





# Observation

- Theorem-- If quadratic probing is used and the table size is prime, then a new element can always be inserted if the table is at least half empty. (see book for more details)



# Example

	Empty Table	After 89	After 18	After 49	After 58	After 69
0				49	49	49
1						
2					58	58
3						69
4						
5						
6						
7						
8			18	18	18	18
9		89	89	89	89	89



# Notes

- If the table is even one more than half full, the insertion could fail (although this is extremely unlikely).
- It is also crucial that the table size be prime.
- If the table size is not prime, the number of alternate locations can be severely reduced.
- Standard deletion cannot be performed in a closed hash table, because the cell might have caused a collision to go past it.
- Closed hash tables require lazy deletion.



# About Lazy Deletion

- Deletion in a hash table is not an easy task. One method to delete an entry is to invent another **special key**, to be placed in any deleted position.
- This special key would indicate that this position is free to receive an insertion when desired but that is **should not** be used to terminate the search for some other item in the table.



# Key-Dependent Increments

- Rather than having the increment depend on the number of probes already made, we can let it be some simple function of the key itself.
- For example, we could truncate the key to a single character and use its code as the increment.



# Random Probing

- Use a pseudorandom number generator to obtain the increment.
- The generator used should be one that always generates the same sequence provided it starts with the same seed.
- This method is excellent in avoiding clustering, but is likely to be slower than the others.



# Double Hashing

- For double hashing, one popular choice is  $f(i) = i * h_2(x)$ .
- We apply a second hash function to  $x$  and probe at a distance  $h_2(x), 2 h_2(x), \dots$ , and so on.
- A poor choice of  $h_2(x)$  would be disastrous.
- The function must never evaluate to zero.
- Make sure all cells can be probed.



# Double Hashing

- For instance, the obvious choice  $h_2(x) = x \bmod 9$  would not help if 99 were inserted into the input in the previous examples.
- A function such as  $h_2(x) = R - (x \bmod R)$ , with  $R$  a prime smaller than  $H\_SIZE$ , will work well.
- One may continue to perform triple hashing, and so on.





# Rehashing

- When the table gets too full, the running time for the operations will deteriorate, specially when there are too many removals intermixed with insertions.
- Solution
  - Build another table that is about twice as big (with associated new hash function).
  - Scan down the entire original hash table, computing the new hash value for each element and inserting it in the new table.



# Example

$$h_2(x) = R - (x \bmod R), \quad R = 7$$

	Empty Table	After 89	After 18	After 49	After 58	After 69
0						69
1						
2						
3					58	58
4						
5						
6				49	49	49
7						
8			18	18	18	18
9		89	89	89	89	89



# Example

- Table size = 7
- Insert 13, 15, 24, and 6.
- $h(x) = x \bmod 7$ .

0	6
1	15
2	
3	24
4	
5	
6	13



# Closed Hash Table, Insert 23

0	6
1	15
2	23
3	24
4	
5	
6	13

- After 23 is inserted, the resulting table will be over 70% full.
- A new table is created with size = 17 since this is the first prime that is twice as large as the old table size.
- The new hash function is then  $h(x) = x \bmod 17$ .



# After Rehashing

The old table is scanned, and elements 6, 15, 23, 24, and 13 are inserted into the new table.

- The running time is  $O(n)$ .
- It is expensive.
- It should not be done so frequently.

0	
1	
2	
3	
4	
5	
6	6
7	23
8	24
9	
10	
11	
12	
13	13
14	
15	15
16	




# Extendible Hashing

- What happens when the amount of data is too large to fit in main memory and must be stored on the disk?
- How can we minimize the disk access?
- Suppose that our data consists of several 6 bit integers.
- The root of the tree contains 4 pointers determined by the leading two bits of the data.



# Example



# Extendible Hashing

- Each leaf has up to  $m = 4$  elements.
- $D$  represents the number of bits used by the root, which is sometimes known as the **directory**.
- The number of entries in the directory is thus  $2^D$   $d_l$  (the number of leading bits that all the elements of some leaf  $l$  have in common).
- $d_l$  will depend on the particular leaf.








# Example, insert 100100

This will go into the third leaf and cause a split.

Now, the leaves are now determined by the first 3 bits.



# Example, insert 000000



# Notes

- It is possible that several directory splits will be required if the elements in a leaf agree in more than  $D+1$  leading bits.
  - For example, 111010, 111011, and 111100 are inserted, the directory size must be increased to 4.
- The possibility of duplicate keys. This algorithm does not work when there are more than  $m$  duplicates.
- It is important for the bits to be fairly random.



# Summary

- Hash tables can be used to implement the **insert** and find operations in constant average time.
- It is especially important to pay attention to details such as load factor when using hash tables.
- It is also important to choose the hash function carefully when the key is not a short string or integer.



# Summary

- For open hashing, the load factor should be close to 1.
- For closed hashing, the load factor should not exceed 0.5, unless this is completely unavoidable.
- Using a hash table, it is not possible to find the minimum element.
- It is not possible to search efficiently for a string unless the exact string is known.



# Summary

- Compilers use hash tables to keep track of declared variables in source code. The data structure is known as a symbol table.
- A hash table is useful for any graph theory problem where the nodes have real names instead of numbers.
- A third common use of hash tables is in programs that play games.
- Another use of hashing is in online spelling checkers.

